

FST Trimming: Ending Dictionary Redundancy in Apertium

Matthew Marting, Kevin Brubeck Unhammer

St. David's School, Kaldera språkteknologi AS
Raleigh, NC., Stavanger, Norway
Ø, unhammer+apertium@mm.st

Abstract

The Free and Open Source rule-based machine translation platform Apertium uses Finite State Transducers (FST's) for analysis, where the output of the analyser is input to a second, *bilingual* FST. The bilingual FST is used to translate analysed tokens (lemmas and tags) from one language to another. We discuss certain problems that arise if the analyser contains entries that do not pass through the bilingual FST. In particular, in trying to avoid “half-translated” tokens, and avoid issues with the interaction between multiwords and tokenisation, language pair developers have created redundant copies of monolingual dictionaries, manually customised to fit their language pair. This redundancy gets in the way of sharing of data and bug fixes to dictionaries between language pairs. It also makes it more complicated to reuse dictionaries outside Apertium (e.g. in spell checkers). We introduce a new tool to *trim* the bad entries from the analyser (using the bilingual FST), creating a new analyser. The tool is made part of Apertium's Ittoolbox package.

Keywords: FST, RBMT, dictionary-redundancy

1. Introduction and background

Apertium (Forcada et al., 2011)¹ is a rule-based machine translation platform, where the data and tools are released under a Free and Open Source license (primarily GNU GPL). Apertium translators use Finite State Transducers (FST's) for morphological analysis, bilingual dictionary lookup and generation of surface forms; most language pairs² created with Apertium use the Ittoolbox FST library for compiling XML dictionaries into binary FST's and for processing text with such FST's. This paper discusses the problem of redundancy in monolingual dictionaries in Apertium, and introduces a new tool to help solve it.

The following sections give some background on how FST's fit into Apertium, as well as the specific capabilities of Ittoolbox FST's; then we delve into the problem of monolingual and bilingual dictionary mismatches that lead to redundant dictionary data, and present our solution.

1.1. FST's in the Apertium pipeline

Translation with Apertium works as a pipeline, where each *module* processes some text and feeds its output as input to the next module. First, a surface form like ‘fishes’ passes through the **analyser** FST module, giving a set of analyses like `fish<n><pl>/fish<vblex><pres>`, or, if it is unknown, simply `*fishes`. Tokenisation is done during analysis, letting the FST decide in a left-right longest match fashion which words are tokens. The compiled analyser technically contains several FST's, each marked for whether they have entries which are tokenised in the regular way (like regular words), or entries that may separate other tokens, like punctuation. Anything that has an analysis is a token, and any other sequence consisting of letters of the alphabet of the analyser is an unknown word token. Anything else can separate tokens.

After analysis, one or more **disambiguation** modules select which of the analyses is the correct one. The **pretransfer** module does some minor formal changes to do with multiwords.

Then a disambiguated analysis like `fish<n><pl>` passes through the **bilingual** FST. Using English to Norwegian as an example, we would get `fisk<n><m><pl>` if the bilingual FST had a matching entry, or simply `@fish<n><pl>` if it was unknown in that dictionary. So a known entry may get changes to both lemma (`fish` to `fisk`) and tags (`<n><pl>` to `<n><m><pl>`) by the bilingual FST. When processing input to the bilingual FST, it is enough that the *prefix* of the tag sequence matches, so a bilingual dictionary writer can specify that `fish<n>` goes to `fisk<n><m>` and not bother with specifying all inflectional tags like number, definiteness, tense, and so on. The tag suffix (here `<pl>`) will simply be carried over.

The output of the bilingual FST is then passed to the **structural transfer** module (which may change word order, ensure determiner agreement, etc.), and finally a **generator** FST which turns analyses like `fisk<n><m><pl>` into forms like ‘fiskar’. Generation is the reverse of analysis; the dictionary which was compiled into a generator for Norwegian can also be used as an analyser for Norwegian, by switching the compilation direction.

A major feature of the Ittoolbox FST package is the support for multiwords and compounds, and the automatic tokenisation of all **lexical units**. A lexical unit may be

- a simple, non-multi-word like the noun ‘fish’,
- a space-separated word like the noun ‘hairy frogfish’, which will be analysed as one token even though it contains a space, but will otherwise have no formal differences from other words,
- a multiword with *inner inflection* like ‘takes out’; this is analysed as `take<vblex><pri><p3><sg>#out` and then, after disambiguation, but before

¹<http://wiki.apertium.org/>

²A *language pair* is a set of resources to translate between a certain set of languages in Apertium, e.g. Basque–Spanish.

bilingual dictionary lookup, turned into `take#out<vblex><pri><p3><sg>` – that is, the uninflected part (called the *lemq*) is moved onto the lemma,

- a token which is actually two words like ‘they’ll’; this is analysed as `prpers<prn><subj><p3><mf><pl>+will<vaux><inf>` and then split after disambiguation, but before bilingual dictionary lookup, into `prpers<prn><subj><p3><mf><pl>` and `will<vaux><inf>`,
- a combination of these three multiword types, like Catalan ‘creure-ho que’, analysed as `creure<vblex><inf>+ho<prn><enc><p3><nt># que` and then moved and split into `creure# que<vblex><inf>` and `ho<prn><enc><p3><nt>` after disambiguation, but before bilingual dictionary lookup.

In addition to the above multiwords, where the whole string is explicitly defined as a path in the analyser FST, we have dynamically analysed compounds which are not defined as single paths in the FST, but still get an analysis during lookup. To mark a word as being able to form a compound with words to the right, we give it the ‘hidden’ tag `<compound-only-L>`, while a word that is able to be a right-side of a compound (or a word on its own) gets the tag `<compound-R>`. These hidden tags are not shown in the analysis output, but used by the FST processor during analysis. If the noun form ‘frog’ is tagged `<compound-only-L>` and ‘fishes’ is tagged `<compound-R>`, the lttoolbox FST processor will analyse ‘frogfishes’ as a single compound token `frog<n><sg>+fish<n><pl>` (unless the string was already in the dictionary as an explicit token) by trying all possible ways to split the word. After disambiguation, but before bilingual dictionary lookup, this compound analysis is split into two tokens, so the full word does not need to be specified in either dictionary. This feature is very useful for e.g. Norwegian, which has very productive compounding.

1.2. The Problem: Redundant data

Ideally, when a monolingual dictionary for, say, English is created, that dictionary would be available for reuse unaltered (or with only bug fixes and additions) in all language pairs where one of the languages is English. Common data files would be factored out of language pairs, avoiding redundancy, giving *data decomposition*. Unfortunately, that has not been the case in Apertium until recently.

If a word is in the analyser, but not in the bilingual translation dictionary, certain difficulties arise. As the example above showed, if ‘fishes’ were unknown to both dictionaries, the output would be `*fishes`, while if it were unknown to only the second, the output of the analyser would be `@fish<n><pl>`, and of the complete translation just `@fish`. Given `*fishes`, a post-editor who knows both languages can immediately see what the original was, while the half-translated `@fish` hides the inflection information

in the source text. Just lemmatising the source text – which removes features like number, definiteness or tense – can skew meaning. For example, if the input were the Norwegian Bokmål ‘Sonny gikk til hundene’, meaning ‘Sonny’s life went to ruin’, a Norwegian Nynorsk translator seeing ‘Sonny gjekk til @hund’ would have to look at the larger context (or the original) to infer that it was the idiomatic meaning, and not ‘Sonny went to (a) dog’. Similarly, in the Bokmål sentence ‘To milliarder ødela livet’, literally ‘Two billions [money] destroyed the life’, the definite form is used to mean ‘his life’; since the lemma of ‘livet’ is ambiguous with the plural, the half-translated ‘To milliarder ødela @liv’ could just as well mean ‘Two billions destroyed lives’.³ But it gets worse: Some languages inflect verbs for *negation*, where the half-translated lemma would hide the fact that the meaning is negative. When translating from Turkish, if you see ‘@öl’, you would definitely need more context or the original to know whether it was e.g. ‘öldürdü’ ‘s/he/it killed’ or ‘öldürmedi’ ‘s/he/it did not kill’ – if the MT is used for gisting, where the user doesn’t know the source language, this becomes particularly troublesome. For single-token cases like this, a workaround for the post-editor is to carry surface form information throughout the pipeline, but as we will see below, this fails with multiwords and compounds, which are heavily used in many Apertium language pairs.

A word not known to the bilingual FST might not have its tags translated (or translated correctly) either. When the transfer module tries to use the half-translated tags to determine agreement, the *context* of the half-translated word may have its meaning skewed as well. For example, some nouns in Norwegian have only plural forms; when translating a singular form from Northern Sámi into Norwegian, the bilingual FST transfers the singular tag to a plural tag for such words. Structural transfer rules insert determiners which are inflected for number; if the noun were half-translated (thus with the singular tag of the input), we would output a determiner with the wrong number, thus the output has two badly translated words instead of untranslated one.

Trying to write transfer rules to deal with half-translated tags also *increases the complexity of transfer rules*. For example, if any noun can be missing its gender, that’s one more exception to all rules that apply gender agreement, as well as any feature that interacts with gender.

Then there are the issues with tokenisation and multiwords. Multiwords are entries in the dictionaries that may consist of what would otherwise be several tokens. As an example, say you have ‘take’ and ‘out’ listed in your English dictionary, and they translate fine in isolation. When translating to Catalan, we want the phrasal verb ‘take out’ to turn into a single word ‘treure’, so we list it as a multiword with *inner inflection* in the English dictionary. This makes any occurrence of forms of ‘take out’ get a single-token multiword analysis, e.g. ‘takes out’ gets the analysis `take<vblex><pri><p3><sg># out`. But then the whole multiword *has* to be in the bilingual dictionary

³Examples from <http://www.bt.no/bergenpuls/litteratur/Variabelt-fra-Nesbo-3081811.html> and <http://www.nettavisen.no/2518369.html>.

if the two words together are to be translated. If another language pair using the same English dictionary has both ‘take’ and ‘out’ in its bilingual dictionary, but not the multiword, the individual words in isolation will be translated, but whenever the whole string ‘take out’ is seen, it will only be lemmatised, not translated. This is both frustrating for the language pair developer, and wasted effort in the case where we don’t *need* the multiword translation.

Assuming we could live with some frustration and transfer complexity, we could try to carry input surface forms throughout the pipeline to at least avoid showing lemmatised forms. But here we run into another problem. Compounds and multiwords that consist of several lemmas are split into two units before transfer, e.g. the French contraction ‘au’ with the analysis `à<pr>+le<det><def><m><sg>` turns into `à<pr>` and `le<det><def><m><sg>`, while the Norwegian compound ‘vasskokaren’ analysed as `vatn<n><nt><sg><ind><cmp>+kokar<n><m><sg><def>` turns into `vatn<n><nt><sg><ind><cmp>` and `kokar<n><m><sg><def>` – compounds may also be split at more than one point. We split the analysis so that we avoid having to list every possible compound in the bilingual dictionary, but we can’t split the form. We don’t even know which part of the form corresponds to which part of the analysis (and of course we would not want to translate half a word). If we attach the full form to the first part and leave the second empty, we run into trouble if only the second part is untranslatable, and vice versa. One hack would be to attach the form to the first analysis and let the second analysis instead have some special symbol signifying that it is the analysis of the form of the previous word, e.g. `vasskokaren/vatn<n><nt><sg><ind><cmp>` and `$1/kokar<n><m><sg><def>`. But now we need the bilingual FST to have a memory of previously seen analyses, which is a step away from being finite-state. It would also mean that other modules which run between the analyser and the bilingual FST in the pipeline have to be careful to avoid changing such sequences,⁴ introducing brittleness into an otherwise robust system.

Due to such issues, most language pairs in Apertium have a separate copy of each monolingual dictionary, manually *trimmed* to match the entries of the bilingual dictionary; so in the example above, if ‘take out’ did not make sense to have in the bilingual dictionary, it would be removed from the copy of the monolingual dictionary. This of course leads to a lot of redundancy and duplicated effort; as an example, there are currently (as of SVN revision 50180) twelve Spanish monolingual dictionaries in stable (SVN trunk) language pairs, with sizes varying from 36,798 lines to 204,447 lines.

The redundancy is not limited to Spanish; in SVN trunk we also find 10 English, 7 Catalan, and 4 French dictionaries. If we include unreleased pairs, these numbers turn to 19, 28, 8 and 16, respectively. In the worst case, if you add some

words to an English dictionary, there are still 27 dictionaries which miss out on your work. The numbers get even worse if we look at potential new language pairs. Given 3 languages, you only need $3 * (3 - 1) = 6$ monolingual dictionaries for all possible pairs (remember that a dictionary provides both an analyser and a generator). But for 4 languages, you need $4 * (4 - 1) = 12$ dictionaries; if we were to create all possible translation pairs of the 34 languages appearing in currently released language pairs, we would need $34 * (34 - 1) = 1122$ monolingual dictionaries, where 34 ought to be enough.




Figure 1: Current number of monodixes with pairs of four languages




Figure 2: Ideal number of monodixes with four languages

The lack of shared monolingual dictionaries also means that other monolingual resources, like disambiguator data, is not shared, since the effort of copying files is less than the effort of letting one module depend on another for so little gain. And it complicates the reuse of Apertium’s extensive (Tyers et al., 2010) set of language resources for other systems: If you want to create a speller for some language supported by Apertium (now possible for Ittoolbox dictionaries via HFST (Pirinen and Tyers, 2012)), you either have to manually merge dictionaries in order to gain from all the work, or (more likely) pick the largest one and hope it’s good enough.

1.3. A Solution: Intersection

However, there is a way around these troubles. Finite state machines can be intersected with one another to produce a new finite state machine. In the case of the Apertium transducers, what we want is to intersect the output (or **right**) side of the full analyser with the input (or **left**) side of the bilingual FST, producing a *trimmed* FST. We call this process *trimming*.

Some recent language pairs in Apertium use the alternative, Free and Open Source FST framework HFST (Linden et al., 2011)⁵. Using HFST, one can create a “pre-fixed” version of the bilingual FST, this is the concatenation of the bilingual FST and the regular expression `.*`,

⁴E.g. the rule based disambiguation system Constraint Grammar allows for moving/removing analyses; and an experimental module for handling discontinuous multiwords also moves words.

⁵<http://hfst.sourceforge.net/>

⁶Partly due to available data in that formalism, partly due to features missing from Ittoolbox like *flag diacritics*.

i.e. match any symbol zero or more times.⁷ Then the command `hfst-compose-intersect` on the analyser and the prefixed FST creates the FST where only those paths of the analyser remain where the right side of the analyser match the left side of the bilingual FST. The prefixing is necessary since, as mentioned above, the bilingual dictionary is underspecified for tag suffixes (typically inflectional tags such as definiteness or tense, as opposed to lemma-identifying tags such as part of speech and noun gender). The HFST solution works, but is missing many of the Apertium-specific features such as different types of tokenisation FST's, and it does not handle the fact that multiwords may split or change format before bilingual dictionary lookup. Also, unlike `lttoolbox`, most of the Apertium dictionaries compiled with HFST represent compounds with an optional transition from the end of the noun to the beginning of the noun dictionary – so if `frog<n>` and `fish<n>` were in the analyser, but `fish<n>` were missing from the bilingual FST, `frog<n>+fish<n>` would remain in the trimmed FST since the prefix `frog<n>.*` matches. In addition, using HFST in language pairs whose data are all in `lttoolbox` format would introduce a new (and rather complex) dependency both for developers, packagers and users who compile from source.

Thus we decided to create a new tool within `lttoolbox`, called `lt-trim`. This tool should trim an analyser using a bilingual FST, creating a trimmed analyser, and handle all the `lttoolbox` multiwords and compounds, as well as letting us retain the special tokenisation features of `lttoolbox`. The end result should be the same as perfect manual trimming. The next section details the implementation of `lt-trim`.⁸

2. Implementation of `lt-trim`

The implementation consists of two main parts: preprocessing the bilingual dictionary, and intersecting it with the analyser.

2.1. Preprocessing the bilingual dictionary

Like monolingual dictionaries, bilingual ones can actually define several FST's, but in this case the input is already tokenised – the distinction is only useful for organising the source, and has no effect on processing. So the first preprocessing step is to take the union of these FST's. This is as simple as creating a new FST F , with epsilon (empty/unlabelled) transitions from F 's initial state, to each initial state in the union, and from each of their final states to F 's final state.

Next, we append loopback transitions to the final state. Like mentioned in section 1.1. above, the bilingual dictionary is underspecified for tags. We want an analyser entry ending in `<n><pl>` to match the bilingual entry ending in `<n>`. Appending loopback transitions to the final state, i.e. `<n>.*`, means the intersection will end up containing `<n><pl>`; we call the bilingual dictionary *prefixed* when

it has the loopback transitions appended. The next section explains the implementation of prefixing.

The final preprocessing step is to give multiwords with inner inflection the same format as in the analyser. As mentioned in section 1.1., the analyser puts tags after the part of the lemma corresponding to the inflected part, with the uninflected part of the multiword lemma coming last.⁹ The bilingual dictionary has the uninflected part before the tags, since it has to allow tag prefixes instead of requiring full tag sequences. Section 2.3. details how we move the uninflected part after the (prefixed) tags in preprocessing the bilingual dictionary.

2.2. Prefixing the bilingual dictionary

The `lttoolbox` FST *alphabets* consist of symbol pairs, each with a left (serves as the input in a transducer) and right (output) symbol. Both the pairs and the symbols themselves, which can be either letters or tags, are identified by an integer. First, we loop through analyser symbol-pairs, collecting identifiers of those where the right-side is a tag. Then for all these tags, we add new symbol-pairs with that tag on both sides to the bilingual FST.

These are used to create loopbacks in the bilingual FST using the function `appendDotStar`. Transitions are created from the final states of the bilingual transducer that loop directly back with each of the identifiers. If we call the set of tag symbol-pairs captured from the analyser T , and the bilingual FST B , the prefixed bilingual FST is BT^* .

2.3. Moving uninflected lemma parts

To turn `take# out<vblex>.*` into `take<vblex>.*# out` and so on, we do a depth-first traversal looking for an occurrence of the `#` symbol. Then we replace the `#`-transition t with one into the new transducer returned by `copyWithTagsFirst(t)`.

This function traverses the FST from the target of t (in the example above, the state before the space), building up two new transducers, *new* and *lemq*. We keep a *SearchState* during traversal, which is a pair of the current source state in the original FST and the last added state in *lemq*. Until we see the first tag, we add all transitions found to *lemq*, and record in the *SearchState* which was the last *lemq* state we saw. In the example above, we would build up a *lemq* transducer containing `out` (with an initial space). Upon seeing the first tag, we start adding transitions from the initial state of *new*,¹⁰ and don't change the last-*lemq*-state of the *SearchState*.

When reaching a final tag state s , we add it and the last seen *lemq* state l to a list of pairs f . After the traversal is done, we loop through each (s, l) in f , creating a temporary copy of *lemq* where the *lemq*-state l has been made the only final state, and adding a `#`-transition from each final tag state s in *new* into that copy. In the example, we would copy the *lemq* into one where the state after the letter `t` were made final, and insert that copy after the final

⁷In this article, we use the regular POSIX format for regular expression; the Xerox format used in HFST-compiled dictionaries differs.

⁸Available in the SVN version of the `lttoolbox` package, see <http://wiki.apertium.org/wiki/Installation>, the code itself is at <https://svn.code.sf.net/p/apertium/svn/trunk/lttoolbox>.

⁹Having the tags “lined up with” or at least close to the inflection they represent eases dictionary writing.

¹⁰We keep state-mapping tables so we can refer to the *new* state corresponding to each original FST state.

state s , the state after the `<vblex>.*`. If, from the original $\#$, there were a *lemq* path that didn't have the same last-*lemq*-state (e.g. `take# out<n>.*` or even `take# out.*`) it would end up in a state that were not final after s , and the path would not give any analyses (such paths are removed by FST *minimisation*). But if a *lemq* path did have the same last state, we would want it included, e.g. `take# part<vblex>.*` to `take<vblex>.*# part`. Thus several *lemq* paths may lead from the $\#$ to the various first tag states, but we only connect those paths which were connected in the original bilingual dictionary. Finally, we return *new* from *copyWithTagsFirst* and look for the next $\#$.

2.4. Intersection

The first method we tried of intersecting FST's consisted of multiplying them. This is a method that is simple to prove correct; however, it was extremely inefficient, requiring a massive amount of memory. First, the states of each transducer were multiplied. This meant that every possible state pair, consisting of a state from the monolingual dictionary and the bilingual dictionary, was assigned a state in the trimmed transducer. Next, each of the transitions were multiplied. As the intersection is only concerned with the output of the monolingual dictionary and the input of the bilingual dictionary, the respective symbols had to match; though very many of them did not, a significant number of matching symbols resulted in transitions to redundant and unreachable states. These were removed with minimisation.

The tool now implements the much more efficient method of intersection, depth-first traversal. Both the monolingual dictionary and the bilingual dictionary are traversed at the same time, in lockstep; only transitions present in both are further followed and added to the trimmed transducer. This depth-first traversal method uses memory and time comparative to regular compilation in *lttoolbox*.¹¹

The process is managed through noting which states are to be processed, the pair of states currently being processed, the next pair of states, and the states which have already been seen. Besides having reachable counterparts in both transducers, a state pair will only be added to the queue if it has not been seen yet.

However, to handle multiwords, a few other things are necessary.

If a $+$ is encountered in the monolingual dictionary (indicating a $+$ -type multiword) the traversal of the bilingual dictionary resumes from its beginning. In addition, in the event that a $\#$ is later encountered, the current position is recorded. The $\#$ -type multiwords alone can be easily handled if the bilingual dictionary is preprocessed to be in the

same format as the monolingual dictionary; the tags must be moved before the $\#$. However, a combination of both $+$ and $\#$ requires that the traversal of the bilingual dictionary return to the state at which the $+$ was first encountered.

We also need some exceptions for epsilon transitions; the idea here is that if we see an epsilon in one transducer, we move across that epsilon without moving in the other transducer, and vice versa. Compound symbols in the analyser are treated similarly: we add the transitions to the trimmed analyser and move forward in the analyser without moving forward in the bilingual FST.

2.5. Using the tool

To the end user (i.e. language pair developer), the tool is a simple command line program with three arguments: the input analyser FST, the input bilingual FST and the output trimmed analyser FST.

3. Ending Dictionary Redundancy

As mentioned in section 1.3., there are already language pairs in Apertium that have moved to a decomposed data model, using the HFST trimming method. At first, the HFST language pairs would also copy dictionaries, even if they were automatically trimmed, just to make them available for the language pair. But over the last year, we have created scripts for our GNU Autotools-based build system that let a language pair have a formal dependency on one or more monolingual data packages¹². There is now an SVN module *languages*¹³ where such monolingual data packages reside, and all of the new HFST-based languages pairs now use such dependencies, which are trimmed automatically, instead of making redundant dictionary copies. Disambiguation data is also fetched from the dependency instead of being redundantly copied.

But most of the released and stable Apertium language pairs use *lttoolbox* and still have dictionary redundancy. With the new *lt-trim* tool, it is finally possible to end the redundancy¹⁴ for the pairs which use *lttoolbox*, with

¹²This both means that source and compiled monolingual files are made available to the make files of the language pair, and that the configure script warns if monolingual data packages are missing. Packagers should be able to use this so that if a user asks their package manager, e.g. `apt-get`, to install the language pair `apertium-foo-bar`, it would automatically install dependencies `apertium-foo` and `apertium-bar` first, and use files from those packages.

¹³<http://wiki.apertium.org/wiki/Languages>

¹⁴One could argue that there is still *cross-lingual* redundancy in the bilingual dictionaries – Apertium by design does not use an interlingua. Instead, the Apertium dictionary crossing tool *crossdics* (Toral et al., 2011) provides ways to extract new translations during development: Given bilingual dictionaries between languages A-B and B-C, it creates a new bilingual dictionary between languages A-C. One argument for not using an interlingua during the translation process is that the dictionary resulting from automatic crossing needs a lot of manual cleaning to root out false friends, unidiomatic translations and other errors – thus an interlingua would have to contain a lot more information than our current bilingual dictionaries in order to automatically disambiguate such issues. It would also require more linguistics knowledge of developers and heighten the entry barrier for new

¹¹When testing on Norwegian Nynorsk-Bokmål, with 64152 entries in the bilingual dictionary and 179369 in the full Bokmål monolingual dictionary, the trimming memory usage is lower than with monolingual compilation (under 10 % on a 4 GB machine), and time is down from 20 to 10 seconds. With English-Catalan, with 27461 entries in the bilingual dictionary and 31784 in the Catalan dictionary and many multiwords, the trimming memory usage is about double of the monolingual compilation (still under 10 %) and time is up from 9 to 13 seconds.

its tokenisation, multiword and compounding features, and without having to make those pairs dependent on a whole other FST framework simply for compilation.

The tool has only recently been released, and there is still much work to do in converting existing language pairs to a decomposed data model. Monolingual dictionaries have to be merged, and the various language pairs may have altered the tag sets in more or less subtle ways that can affect disambiguation, transfer and other parts of the pipeline.

To give an example, merging the Norwegian Bokmål dictionaries of the language pairs Northern Sámi-Bokmål and Nynorsk-Bokmål (including changes to transfer and to the other involved dictionaries) took about three hours of work. However, this kind of merge work happened often in the past anyway when major changes happened to either dictionary; from now on it can be a one-time job. Future additions and changes to the common apertium-nob module will benefit both language pairs.

Any new languages added to Apertium can immediately reap the benefits of the tool, without this manual merge work; this goes for many of the in-development pairs too (e.g. the English-Norwegian language pair now depends on the Bokmål and Nynorsk monolingual package).

4. Conclusion

In this article, we have presented a new tool to trim monolingual dictionaries to fit the data in Apertium language pairs. The tool has already been implemented and used in several language pairs, starting the process to end monolingual dictionary redundancy.

In future, we plan to research how to deal with compounds when trimming HFST dictionaries, as well as further merging monolingual dictionaries.

Acknowledgements

Part of the development was funded by the Google Code-In¹⁵ programme.

5. References

- Mikel L. Forcada, Mireia Ginestí-Rosell, Jacob Nordfalk, Jim O'Regan, Sergio Ortiz-Rojas, Juan Antonio Pérez-Ortiz, Felipe Sánchez-Martínez, Gema Ramírez-Sánchez, and Francis M. Tyers. 2011. Apertium: a free/open-source platform for rule-based machine translation. *Machine Translation*, 25(2):127–144.
- Krister Linden, Miikka Silfverberg, Erik Axelsson, Sam Hardwick, and Tommi Pirinen, 2011. *HFST—Framework for Compiling and Applying Morphologies*, volume Vol. 100 of *Communications in Computer and Information Science*, pages 67–85.
- Tommi A Pirinen and Francis M Tyers. 2012. Compiling apertium morphological dictionaries with hfst and using them in hfst applications. In G. De Pauw, G-M de Schryver, M.L. Forcada, K. Sarasola, F.M. Tyers, and P.W. Wagacha, editors, *[SALTMIL 2012] Workshop on Language Technology for Normalisation of Less-Resourced Languages*, pages 25–28, May.
- Antonio Toral, Mireia Ginestí-Rosell, and Francis Tyers. 2011. An Italian to Catalan RBMT system reusing data from existing language pairs. In F. Sánchez-Martínez and J.A. Pérez-Ortiz, editors, *Proceedings of the Second International Workshop on Free/Open-Source Rule-Based Machine Translation*, pages 77–81, Barcelona, Spain, January.
- Francis M. Tyers, Felipe Sánchez-Martínez, Sergio Ortiz-Rojas, and Mikel L. Forcada. 2010. Free/Open-Source Resources in the Apertium Platform for Machine Translation Research and Development. *Prague Bull. Math. Linguistics*, 93:67–76.

contributors.

¹⁵<https://code.google.com/gci/>